



Co-funded by  
the European Union



**EuroHPC**  
Joint Undertaking

# Multi**scale**

**EuroHPC JU Centre of Excellence**

*EESSI test suite*


**10<sup>th</sup> EasyBuild User Meeting@ Juelich**

Thu 27 March 2025

Caspar van Leeuwen (SURF)  
Lara Peeters (Ugent)  
Sam Moors (VUB)



# The EESSI test suite



Goal of the EESSI test suite

- ◎ To test the functionality and performance of the EESSI software stack on a wide range of systems

The opportunity

- ◎ Same software (modules) on every system

The challenge

- ◎ Every system is different! Need tests that are portable

# Writing portable tests is challenging...

- ◎ EESSI test suite is based on ReFrame
- ◎ ReFrame tests are *typically* very system specific, example attributes:
  - `valid_systems`: on which systems this test can/should run
  - `num_cpus_per_task`, `num_tasks`, `num_gpus_per_node`: typically chosen to match the system
  - And many more ...
- ◎ ReFrame offers *amazing* fine-grained control, but at the cost of portability

# How we make EESSI tests portable

- ◎ All system-specific information goes into ReFrame config file
- ◎ Make the test do something sensible *based on the config file*, examples:
  - Launch one rank per available (physical) CPU core / numa node / socket / GPU
  - Skip a test if the system has insufficient memory to run it
  - Skip a test that requires GPUs if the system doesn't have any

N.B. Tests  $\neq$  benchmarks! These portable tests are *not* guaranteed to get the best performance from your system for a particular use case, they are meant to spot performance changes.

# MPI4PY example

```
@rfm.simple_test
class EESSI_MPI4PY(rfm.RunOnlyRegressionTest, EESSI_Mixin):
    device_type = DEVICE_TYPES[CPU]
    compute_unit = COMPUTE_UNIT[CPU]

    module_name = parameter(find_modules('mpi4py'))

    n_iterations = variable(int, value=1000)
    n_warmup = variable(int, value=100)

    executable = 'python3'
    executable_opts = ['mpi4py_reduce.py', '--n_iter', f'{n_iterations}', '--n_warmup', f'{n_warmup}']

    time_limit = '5m00s'

    bench_name = 'mpi4pi'
    bench_name_ci = 'mpi4pi'

    readonly_files = ['mpi4py_reduce.py']

    def required_mem_per_node(self):
        return self.num_tasks_per_node * 100 + 250

    @sanity_function
    ...
    @performance_function('s')
    ...
```

Requires 'CPU' feature

Launch one task per core

Create tests for all modules called mpi4py/<something>


Automatically tags a test instance with this bench\_name with 'CI' tag

Request sufficient memory, and skip if nodes don't have enough

# How we use the EESSI test suite

- ◎ We run the test suite periodically (daily/weekly) on about 5 different systems
- ◎ We run selected (single node) tests when building new software for EESSI (before deployment)
- ◎ Some of us also use the test suite to test local module stacks!

System=snellius:rome • Software stack=production • module\_name=TensorFlow/2.13.0-foss-2023a • device\_type=cpu



# Writing an EESSI test suite configuration

Goal: For everyone to have run the EESSI test suite on your HPC cluster (or laptop) by the end of EUM'25!

- ◎ Step 1: install ReFrame & the EESSI test suite
- ◎ Step 2: create a ReFrame configuration file
- ◎ Step 3: run `reframe --list -t CI`
- ◎ Step 4: run `reframe --dryrun -t CI -n /<somehash>`
- ◎ Step 5: run `reframe --run -t CI -n /<somehash>`



# Find presentation online

- ◎ Detailed steps on subsequent slides
- ◎ Quickest way: copy-paste from slides at <https://github.com/casparvl/EUM25>

# Step 1: Install ReFrame & EESSI test suite

```
module purge # Use system python
python3 -m venv $HOME/eessi_testsuite/eessi_testsuite_venv
source $HOME/eessi_testsuite/eessi_testsuite_venv/bin/activate
pip install reframe-hpc
pip install eessi-testsuite
# Check we can use things from ReFrame's hpctestlib
python3 -c 'import hpctestlib.sciapps.gromacs'
# Check we can use things from the EESSI testsuite
python3 -c 'import eessi.testsuite.eessi_mixin'
```

## Step 2a: create ReFrame config file

```
cd $HOME/eessi_testsuite/  
wget https://raw.githubusercontent.com/casparvl/test-suite/refs/heads/update\_example\_settings/config/settings\_example.py  
export RFM_CONFIG_FILES=$HOME/eessi_testsuite/settings_example.py  
export RFM_PREFIX=$HOME/eessi_testsuite/reframe_runs  
export  
RFM_CHECK_SEARCH_PATH=$HOME/eessi_testsuite/eessi_testsuite_venv/lib  
/python3.9/site-packages/eessi/testsuite/tests/  
export RFM_CHECK_SEARCH_RECURSIVE=1
```

## Step 2b: create ReFrame config file

Now, modify `settings_example.py` to match your system

- ⦿ Define a `stagedir` on a shared filesystem
- ⦿ Select the matching scheduler [https://reframe-hpc.readthedocs.io/en/stable/config\\_reference.html#config.systems.partitions.scheduler](https://reframe-hpc.readthedocs.io/en/stable/config_reference.html#config.systems.partitions.scheduler) (set `local` if you are doing this on your laptop)
- ⦿ Select the matching parallel launcher [https://reframe-hpc.readthedocs.io/en/stable/config\\_reference.html#config.systems.partitions.launcher](https://reframe-hpc.readthedocs.io/en/stable/config_reference.html#config.systems.partitions.launcher) (`mpirun` *should* work for everyone, but you can use e.g. `srun` )
- ⦿ Modify the `access` field to define arguments to be passed to the scheduler, etc. It should define a homogeneous set of nodes

## Step 2c: create ReFrame config file

Now, modify `settings_example.py` to match your system

- ② Under `resources` set the flag that should be passed to your scheduler to define required memory per node and pass `{size}` as argument
  - Slurm users: `--mem={size}`
  - Local spawner: `--whatever={size}` (unused)
- ② Define the max available memory per node under the `EXTRAS.MEM_PER_NODE` item (in MiB).
  - SLURM users: check `scontrol show node <nodename>` for the RealMemory on your nodes.
  - Local spawner: put anything (unused)

## Step 2d: create ReFrame config file

Now, modify `settings_example.py` to match your system

- ◎ Under `features` specify what FEATURES (CPU/GPU) and SCALES your system support
  - CPU partition: `'features': [FEATURES.CPU],`
  - GPU partition where you don't want to run CPU-only tests: `'features': [FEATURES.GPU],`
  - GPU partition where you also want to run CPU-only tests: `'features': [FEATURES.CPU, FEATURES.GPU],`
  - To run all scales (up to 16 nodes): `'features': [FEATURES.XYZ] + list(SCALES.keys())`
  - To run only single (full) node (e.g. local laptop): `'features': [FEATURES.XYZ] + [key for key, value in SCALES.items() if value.get("num_nodes") == 1]`
- ◎ GPU partitions only: under `extras` define `EXTRAS.GPU_VENDOR: GPU_VENDORS.NVIDIA`

## Step 3: run reframe --list -t CI

Run `reframe --list -t CI`

- ◎ You may get things like “WARNING: skipping test 'EESSI\_TensorFlow': the following parameters are undefined: module\_name”. That’s ok, it simply means you don’t have the software(module) needed to run this test

## Step 3: run `reframe --list -t CI`

Run `reframe --list -t CI`

- ◎ If you get “WARNING: failed to retrieve remote processor info: command 'sbatch rfm-detect-job.sh' failed with exit code 1:”, ReFrame’s automatic CPU detection failed.
  - Check the ReFrame log (“Log file(s) saved in ‘/path/to/log’”)
  - You might be missing `access` arguments
  - If it keeps failing, you could try ‘manually’ running `reframe --detect-host-topology` on the relevant node <https://www.eessi.io/docs/test-suite/ReFrame-configuration-file/#create-topology-file> . Then copy to `~/.reframe/topology/<system>-<partition>/processor.json`



## Step 3: run `reframe --list -t CI`

Run `reframe --list -t CI`

- ◎ You'll need to have at least one module available for which we have a test ◎
  - If you don't, simply install e.g. a CPU version of OSU-MicroBenchmarks with EasyBuild

◎ Expected output:

...

- EESSI\_TensorFlow %scale=2\_nodes %module\_name=TensorFlow/2.13.0-foss-2023a %device\_type=cpu /cbc475c5
- EESSI\_TensorFlow %scale=1\_node %module\_name=TensorFlow/2.13.0-foss-2023a %device\_type=cpu /9864d0f5



Test hash

## Step 4: run reframe --dryrun -t CI


Run `reframe --dryrun -t CI -n /<testhash>` to just run an individual test as an example

- ⦿ Check the jobscript ReFrame will generate & submit in  
`<stagedir>/<system_name>/<partition_name>/default/<testname_testhash>/rfm_job.sh`
- ⦿ If you have issues, that job script is (probably) your first place to look!

## Step 5: run reframe --run -t CI

Run `reframe --run -t CI -n /<testhash>` to just run an individual test as an example

# Summary



- ① Writing the ReFrame config requires some knowledge specific to the EESSI test suite, but is documented <https://www.eessi.io/docs/test-suite/ReFrame-configuration-file/>
- ① Apart from the ReFrame config, the EESSI test suite is ‘plug-and-play’!
- ① Number of supported applications is could be bigger – open to new contributions (see <https://www.eessi.io/docs/test-suite/writing-portable-tests/>)
- ① EESSI test suite is application-focused: no replacement for system-specific (e.g. scheduler or filesystem) tests
- ① Credit to ReFrame devs: EESSI test suite is possible because they spent time on our bug reports & feature requests 😊

# MultiXscale

Web page: [multixscale.eu](http://multixscale.eu)

Facebook: [MultiXscale](https://www.facebook.com/MultiXscale)

X: [@MultiXscale](https://twitter.com/MultiXscale)

LinkedIn: [multixscale](https://www.linkedin.com/company/multixscale)

YouTube: [@MultiXscale](https://www.youtube.com/channel/UCMultiXscale)



Co-funded by  
the European Union



**EuroHPC**  
Joint Undertaking



UNIVERSITAT DE  
BARCELONA



Universität  
Stuttgart



**SORBONNE  
UNIVERSITÉ**



Université  
de Toulouse



Consiglio Nazionale  
delle Ricerche



MAX-PLANCK-GESELLSCHAFT



Funded by the European Union. This work has received funding from the European High Performance Computing Joint Undertaking (JU) and countries participating in the project under grant agreement No 101093169.